New York voter file codebook

December 2022 statewide snapshot

File: NY_2022_12.parquet | **Format:** Apache Parquet | **Size:** 1,088.7 MB. Read with R arrow or Python pyarrow; see the README for examples.

Coverage: all 62 New York counties, including New York City. The source file is named AllNYSVoters_20221227.txt. December 27, 2022 is the date in that source filename; this is a single snapshot, not a panel.

Contents: 20,919,063 rows, 114 columns and 19,937,571 distinct state_vid values. IDs are never null or blank. The same ID can occur under multiple registration statuses. The file has 981,492 more rows than distinct IDs, so a row is not necessarily a distinct voter.

This release uses the validated coordinate-corrected version from October 2026. Only the latitude/longitude column assignment differs from the earlier standardized file; other values and row order are preserved.

Registration status

Status	Rows	Distinct IDs	Interpretation
ACTIVE	12,168,279	12,168,275	Registration marked active by the pipeline.
INACTIVE	972,235	972,235	Registration marked inactive; this does not mean a person never votes.
PURGED	7,694,139	7,694,139	Purged registration record retained in the source; exclude from current-registration counts.
INELIGIBLE	84,410	84,410	Pipeline category for records not currently effective, including preregistrations. Not a legal assessment of the person.

Distinct IDs by status overlap and should not be summed to obtain the number of unique IDs in the whole file.

Before using the data

• Count distinct state_vid values within a clearly defined status group. Inspect repeated records before selecting one row per ID for attribute analysis.

• The file includes 74 voted_YYYYMMDD indicators spanning 1992-2022. These record participation, never vote choice. Historical coverage and zero/missing conventions require care; see pages 5-6.

• Race is unavailable. Some birth and registration dates are implausible. Mailing fields are only partly parsed. Geocodes vary in precision; details follow.

• These are administrative registration records, not a sample or census of residents. Historical, inactive and purged records should not be treated as the current eligible population.

Class use only

Provided to Mark Hansen and his Columbia Journalism School data class. Do not redistribute the data or identifiable record-level extracts outside the class without prior written permission from Kasey Rhee and Andrew C. W. Myers. Keep the Drive folder restricted to authorized class participants. See the accompanying DO NOT SHARE notice.

Identifiers and voter characteristics

Types are the Parquet file's logical storage types. Dates are stored as date32 and normally displayed as YYYY-MM-DD. Keep IDs and ZIP codes as text. Names and other text were normalized by the pipeline.

Variable	Storage	Meaning and cautions
file_id	Text	Internal source-file identifier. Constant NY-2343 in this release.
snapshot_id	Number	Internal snapshot number. Constant 40; not a voter identifier or calendar date.
snapshot_sub_id	Number	Internal component-file number. Entirely missing in this file.
state	Text	State abbreviation. Constant NY.
state_vid	Text	Standardized New York voter-record ID, derived from SBOEID. Keep as text. It is not unique across all rows; the same ID can occur under different statuses. No IDs are null or blank.
reg_date	Date	Registration date carried from the source record. Not the snapshot date and not necessarily the person's first-ever registration. Implausible values remain: observed range 1820-05-18 to 6200-07-07.
updated_date	Date	Standard schema field for a record update date. Entirely missing here.
last_contact_date	Date	Standard schema field for the last contact date. Entirely missing here.
status	Text	Standardized registration status: ACTIVE, INACTIVE, PURGED or INELIGIBLE. See page 1 for counts and interpretation. No original status column is included.
name_first	Text	First name after basic text normalization, including uppercasing. Some entries are blank.
name_middle	Text	Middle name or initial after normalization. Blank entries are common.
name_last	Text	Last name after normalization. Some entries are blank.
name_suffix	Text	Name suffix after normalization, such as JR or III. Usually blank.
sex	Text	Harmonized source sex/gender field. FEMALE or MALE; MISSING DATA for U or blank; OTHER for source X or I. Do not infer a person's gender identity from this field.
race	Text	Constant NOT REPORTED. This source supplies no usable race data; it is not an inferred race measure.
party	Text	Registered party enrollment after harmonization: DEM, REP, NONPARTISAN or OTHER. This is not a measure of vote choice. See the source-code mapping on page 3.
birth_year	Number	Year extracted from source date of birth. Implausible years remain; observed range 1827 to 9182. Validate before use.
birth_date	Date	Parsed source date of birth. Implausible values remain; observed range 1827-08-17 to 9182-09-11. Do not publish identifiable dates of birth.
age	Number	Approximate age at the snapshot, generally 2022 minus birth_year when that field is available. It does not account for the birthday. Values outside 18-120 are missing; 106,574 rows have missing age.

Addresses and source categories

Variable	Storage	Meaning and cautions
voterfile_county	Text	County name mapped from the source county code. All 62 New York counties are represented. Uppercase names, not county FIPS codes.
reg_address_street	Text	Registration street address assembled from source address parts, including apartment information when available. 130,427 rows contain a blank string.
reg_address_zip	Text	Registration ZIP code stored as text. The pipeline strips leading zeros. This is not guaranteed to be a validated five-digit ZIP code.
reg_address_city	Text	Source registration-address city or locality. Uppercase; not a unique geographic identifier.
mail_address_street	Text	Concatenation of up to four mailing-address lines. It can include city, state or postal text within this field; it is not a reliably parsed street-only address. Usually blank.
mail_address_zip	Text	Standard mailing ZIP field. Entirely missing in this release.
mail_address_city	Text	Populated from the raw TOWNCITY field by the importer. Treat as a source town/city value or code, not a validated mailing city.
mail_address_state	Text	The importer sets this to NY for every row. It does not verify the actual mailing state; mailing lines can contain other destinations.
unstandardized_party	Text	Party code before cross-state party harmonization, after basic text cleaning. Twelve observed codes are listed below. This is not a verbatim raw-file extract.
unstandardized_sex	Text	Sex/gender code before harmonization, after basic text cleaning. Observed: F, M, U, X, I and one blank string.
unstandardized_race	Text	Pre-harmonization race field. Entirely missing here.

Party code mapping in this file

Source code	Standardized party	Rows
DEM	DEM	10,208,223
REP	REP	4,831,077
BLK	NONPARTISAN	4,752,562
CON, GRE, IND, LBT, OTH, REF, SAM, WEP, WOR	OTHER	1,127,201

BLK denotes no party enrollment in this mapping. IND is retained as a party code and maps to OTHER; it must not be treated as the no-party category. OTHER combines several source parties, so use unstandardized_party when those distinctions matter.

Missing and blank values

Missing information appears as Parquet nulls (usually NA in R), empty strings, or explicit category labels such as MISSING DATA and NOT REPORTED. These are different representations. snapshot_sub_id, updated_date, last_contact_date, mail_address_zip and unstandardized_race are entirely null. A non-null field is not necessarily valid or informative.

Geocoding and county

Variable	Storage	Meaning and cautions
geocoded	Integer	Processing flag from the address-geocoding table: 1 means processed; 51 rows are null. A value of 1 does not guarantee a usable or precise match. Check coordinates, score and type.
geocoded_lat	Number	Registration-address latitude in decimal degrees, WGS84. Coordinate-corrected release: this column contains latitude. 327 rows lack coordinates.
geocoded_lon	Number	Registration-address longitude in decimal degrees, WGS84. Negative values are west of Greenwich. Coordinate-corrected release: this column contains longitude.
geocoded_score	Number	ArcGIS address-match score from 0 to 100. Higher means a stronger textual match; it is not a probability or a measure of positional error.
geocoded_type	Text	Match level. Main values: POINTADDRESS, STREETADDRESS, STREETNAME and SUBADDRESS. Other observed values: POSTAL, STREETADDRESSEXT, POI, POSTALLOC, STREETINT and LOCALITY, plus blank/null.
geocoded_county	Text	County reported by the geocoder for the registration address. May disagree with voterfile_county or be missing. County-like suffixes are removed.
geocoded_puma_2000	Integer	2000-vintage Public Use Microdata Area assigned from the registration geocode. Integer storage drops leading zeros; format as a five-digit text code for joins.
geocoded_puma_2010	Integer	2010-vintage Public Use Microdata Area assigned from the registration geocode. Use matching-vintage Census geography and a five-digit PUMA code.
geocoded_puma_2020	Integer	2020-vintage Public Use Microdata Area assigned from the registration geocode. Use matching-vintage Census geography and a five-digit PUMA code.
county	Text	Combined county field: uses voterfile_county when populated, otherwise the geocoded county. The source county is populated on every row here, so this field uses the source county.

Using the geographic fields

The coordinates describe registration addresses, not mailing addresses. POINTADDRESS and SUBADDRESS are generally more specific than interpolated STREETADDRESS, street-only or postal matches. There are 327 rows without finite coordinate pairs; 48 finite pairs lie outside broad New York bounds (latitude 40-46 and longitude -80 to -71). Check matches before mapping or assigning neighborhoods.

PUMA codes are not nationally unique. Join on New York state FIPS 36 plus a five-digit PUMA code, using the same Census vintage as the field. A PUMA is a statistical area, not necessarily a neighborhood; boundaries differ across vintages.

Read and summarize with R and Arrow

The file exceeds a spreadsheet worksheet's capacity. Arrow can query selected columns without loading the whole file into memory. This example counts distinct active IDs by source-preferred county; it does not count all rows.

```
library(arrow)
library(dplyr)
ds <- open_dataset("NY_2022_12.parquet")
ds |> filter(status == "ACTIVE") |>
  group_by(county) |>
  summarise(voters = n_distinct(state_vid)) |> collect()
```

Voting history

The 74 added columns are named voted_YYYYMMDD, where YYYYMMDD is the election date encoded by the pipeline. For example, voted_20201103 refers to November 3, 2020 and voted_20221108 to November 8, 2022. All are nullable 32-bit integers in Parquet.

Value	Meaning
`1`	Participation is recorded for that voter ID and election.
`0`	No participation is recorded in the processed history. Some zeros are inferred by the pipeline from absence in a vote list or by adding eligible roster IDs with no history. Zero does not establish that a person was registered and eligible at the historical election.
Null / NA	No matching December 2022 history row was available for this voter ID. This is not a zero. Preserve nulls when analyzing the file.

How the history was attached

The matching December 2022 standardized history file contains 16,492,459 unique IDs. It was left-joined to the coordinate-corrected registration snapshot by state_vid. Each history ID is unique, so the join preserves the original 20,919,063 registration rows and their order. All 40 registration fields remain in place; 74 election columns are appended.

Registration status	Rows with history	Rows without history
`ACTIVE`	12,168,279	0
`INACTIVE`	972,235	0
`PURGED`	4,333,321	3,360,818
`INELIGIBLE`	23	84,387

What students should check

Historical participation is available only for IDs retained in this December 2022 roster, including its inactive and purged records. It omits people who voted earlier but are absent from this snapshot. These counts are not official statewide historical turnout totals. Coverage is especially sparse in the oldest columns: the history source records only 65 positive entries for November 1992.

Before calculating a turnout rate, define the election's relevant registered cohort, account for age and registration timing, and handle missing histories explicitly. An ID may appear in multiple registration rows, so count distinct IDs. Current party, age, address and status describe this snapshot, not necessarily the earlier election.

Some source history entries lack a year or complete date; the importer reconstructs them using chronological position and election-date lookup tables. Not every date has been independently audited. The file reports participation only; candidate choice and voting method are not included.

Election indicator columns

Every column below uses the same 1 / 0 / null conventions on page 5. Names encode YYYYMMDD. Read each column vertically in chronological order, then continue with the next column.

Election indicator	Election indicator	Election indicator
voted_19921103	voted_20040302	voted_20140909
voted_19930914	voted_20040914	voted_20141104
voted_19931102	voted_20041102	voted_20150910
voted_19940913	voted_20050913	voted_20151103
voted_19941108	voted_20051108	voted_20160419
voted_19950912	voted_20060912	voted_20160628
voted_19951107	voted_20061107	voted_20160913
voted_19960307	voted_20070918	voted_20161108
voted_19960910	voted_20071106	voted_20170912
voted_19961105	voted_20080205	voted_20171107
voted_19970909	voted_20080909	voted_20180424
voted_19971104	voted_20081104	voted_20180626
voted_19980915	voted_20090908	voted_20180913
voted_19981103	voted_20091103	voted_20181106
voted_19990914	voted_20100914	voted_20190625
voted_19991102	voted_20101102	voted_20191105
voted_20000307	voted_20110913	voted_20200428
voted_20000912	voted_20111108	voted_20200623
voted_20001107	voted_20120424	voted_20201103
voted_20010911	voted_20120626	voted_20210622
voted_20011106	voted_20120913	voted_20211102
voted_20020910	voted_20121106	voted_20220628
voted_20021105	voted_20130910	voted_20220823
voted_20030909	voted_20131105	voted_20221108
voted_20031104	voted_20140624	

Example 2022 participation summary

This counts distinct active December 2022 IDs with recorded participation in the November 2022 general election, by county. It is a snapshot-based descriptive count, not an official turnout total or a turnout rate.

```
ds |> filter(status == "ACTIVE", voted_20221108 == 1L) |>
  group_by(county) |>
  summarise(recorded_voters = n_distinct(state_vid)) |> collect()
```